


A very quick introduction to  
Auto-encoding variational Bayes and  
the reparameterization trick

J. Bill

February 12, 2019

# The problem with inference in and learning of generative models

Consider a generative model:



We want to perform

- ▶ Inference:  $p_{\theta}(z | x) \propto p_{\theta}(x | z) \cdot p_{\theta}(z)$
- ▶ Learning:  $\Delta\theta = \eta \cdot \nabla_{\theta} \mathbb{E}_{p^*(x)} [\log p_{\theta}(x)]$

Problem:

For expressive model classes  $p_{\theta}(x, z)$ , typically either

- ▶ inference is intractable (e.g., neural nets), or
- ▶ learning becomes computationally expensive (e.g., Boltzmann machines).

## Idea: Separate inference model and learning model

Define a variational posterior,

$$q_{\phi}(z | x) \approx p_{\theta}(z | x) \ ,$$

that makes inference tractable while keep learning the model,

$$\nabla_{\theta} \log p_{\theta}(x) \ ,$$

easy.

We have to keep  $q_{\phi}(z | x)$  and  $p_{\theta}(z | x)$  “consistent” with each other while learning their separate parameters  $\phi$  and  $\theta$ .

This will be formalized in the objective function.

# The evidence lower bound (ELBO) of the data log-likelihood

Three identical ways of writing the ELBO that grant different insights:

$$\mathcal{L}(\theta, \phi; x) := \log p_{\theta}(x) - D_{\text{KL}}[q_{\phi}(z|x) || p_{\theta}(z|x)] \quad (1)$$

$\rightarrow \nabla_{\phi}$  : Keep the variational posterior close the correct one.

$$= \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x, z)] + \mathbb{H}[q_{\phi}(z|x)] \quad (2)$$

$\rightarrow \nabla_{\theta}$  : Improve the generative model (alternate with (1) to  $\uparrow \log p_{\theta}(x)$ )

$$= \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)] - D_{\text{KL}}[q_{\phi}(z|x) || p_{\theta}(z)] \quad (3)$$

$\rightarrow$  Auto-encode the input while keeping the avg. posterior close to the prior.

The decomposition holds for any choice of  $q_{\phi}(z|x)$ !

We can choose  $p_{\theta}(x, z)$  and  $q_{\phi}(z|x)$  such that the gradients  $\nabla_{\theta} p_{\theta}(x, z)$  and  $\nabla_{\phi} q_{\phi}(z|x)$  are easy to compute while evading the complex posterior of  $p_{\theta}(x, z)$ .

## A technical problem: High variance gradient estimates

Gradient ascent on the ELBO always involves terms

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(z|x)} [f(z)] \quad .$$

Yet, we can bring the gradient into the form,

$$\mathbb{E}_{q_{\phi}(z|x)} [f(z) \cdot \nabla_{\phi} \log q_{\phi}(z|x)] \quad .$$

to make it accessible to sample-based estimation. Hooray!

Hooray..?

It turns out that low probability samples from  $q_{\phi}(z|x)$  can have a strong contribution via  $\nabla_{\phi} \log q_{\phi}(z|x)$ . This makes the total gradient estimate very “noisy”!

To noisy, actually, for practical purposes ☹

## Solution: The reparameterization trick

Choose  $q_\phi(z | x)$  such that

$$z = g_\phi(x, \epsilon) \text{ with } \epsilon \sim p(\epsilon) \leftarrow \text{simple distribution.}$$

Then:

$$\nabla_\phi \mathbb{E}_{q_\phi(z | x)} [f(z)] = \nabla_\phi \mathbb{E}_{p(\epsilon)} [f(g_\phi(x, \epsilon))] = \mathbb{E}_{p(\epsilon)} [\nabla_\phi f(g_\phi(x, \epsilon))].$$

Now,  $\mathbb{E}_{p(\epsilon)} [\cdot]$  is easy to sample from while covering the  $q_\phi(z | x)$  space.  
And  $\nabla_\phi f(g_\phi(x, \epsilon))$  is a “backprop-style” chain rule.

It is found, that this estimator has much lower variance, i.e., is less “noisy”.

## AEVB in practice

In practice, we consider the form (3) of the ELBO for learning:

$$\mathcal{L}(\theta, \phi; x) = \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)] - D_{\text{KL}} [q_{\phi}(z|x) || p_{\theta}(z)]$$

We choose to fix the prior,  $p_{\theta}(z) \equiv p(z)$ , thereby turning it into a target distribution for  $q_{\phi}(z|x)$ . Further, we can often choose  $p(z)$  and  $q_{\phi}(z|x)$  such that

$$\nabla_{\phi} D_{\text{KL}} [q_{\phi}(z|x) || p(z)]$$

can be calculated analytically ( $\rightarrow$  no sampling error).

Then, we only apply the reparameterization trick to  $\mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)]:$

Sample-based part of the gradient estimation

$$\mathbb{E}_{p(\epsilon)} [\nabla_{\theta/\phi} \log p_{\theta}(x | g_{\phi}(x, \epsilon))]$$